

RTS Project 1

Part I: Linux Kernel Building

Part II: Linux Scheduling Policy Testing

Advisor: Prof. Tei-Wei Kuo


TA: Han-Yi Lin

Outline

- ▶ What is “Kernel”
- ▶ Environment Setup
- ▶ Build Linux Kernel
- ▶ Test Linux Scheduling Policy
- ▶ Project Requirements
- ▶ Submission Rules

What is “Kernel”?

- ▶ The kernel^[1] is a fundamental part of a modern computer's operating system.
- ▶ The kernel's primary functions are to
 - ▶ Manage the computer's hardware and resources
 - ▶ E.g., CPU, main memory, I/O devices, and so on.
 - ▶ Allow applications to run and use these resources

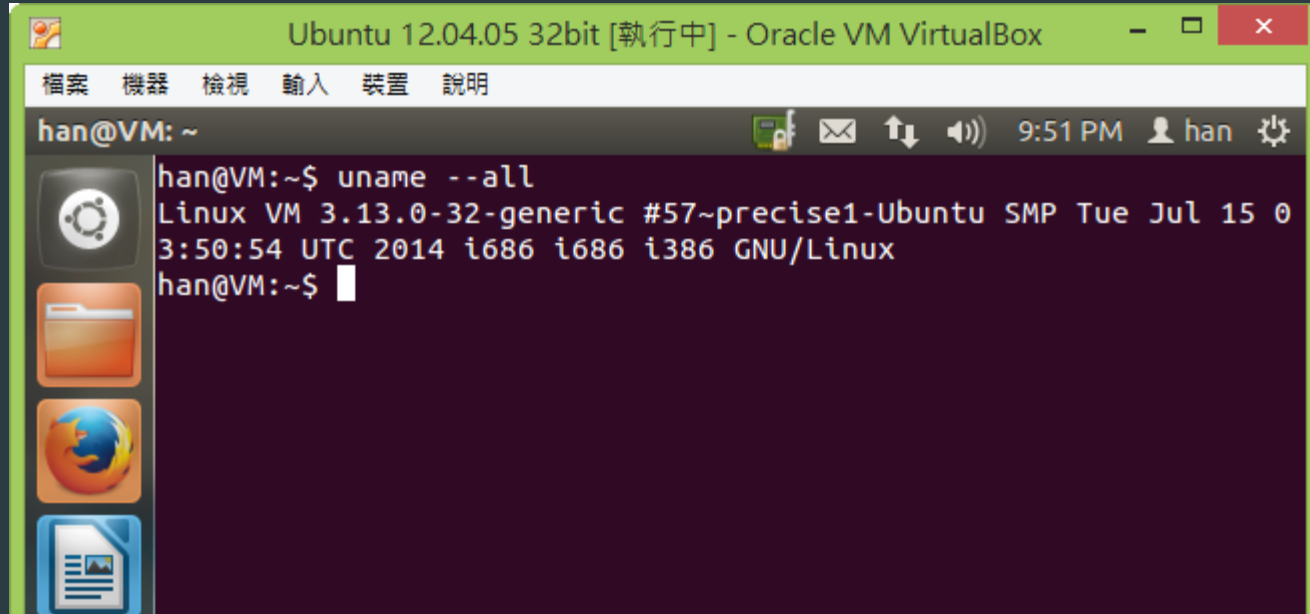


Environment Setup

- ▶ Oracle VM VirtualBox^[2]
 - ▶ Download link:
<https://www.virtualbox.org/wiki/Downloads>
- ▶ Ubuntu 12.04.5 32bits LTS^[3]
 - ▶ Download link:
<http://tw.archive.ubuntu.com/ubuntu-cd/12.04.5/ubuntu-12.04.5-desktop-i386.iso>
- ▶ Install the Ubuntu 12.04.5 on the VirtualBox

Build Linux Kernel (1/5)

- After the installation, please login Ubuntu and open a terminal to start building your Linux kernel^[4]



The screenshot shows a VirtualBox window titled "Ubuntu 12.04.05 32bit [執行中] - Oracle VM VirtualBox". Inside the window, a terminal window is open with the prompt "han@VM: ~". The terminal shows the command "uname --all" being executed, resulting in the output: "Linux VM 3.13.0-32-generic #57~precise1-Ubuntu SMP Tue Jul 15 03:50:54 UTC 2014 i686 i686 i386 GNU/Linux". The terminal window has a dark purple background and a light blue border. The VirtualBox window has a green title bar and a menu bar with options: 檔案, 機器, 檢視, 輸入, 裝置, 說明. The status bar at the bottom of the terminal window shows the time "9:51 PM" and the user "han".

```
han@VM: ~  
han@VM:~$ uname --all  
Linux VM 3.13.0-32-generic #57~precise1-Ubuntu SMP Tue Jul 15 0  
3:50:54 UTC 2014 i686 i686 i386 GNU/Linux  
han@VM:~$
```

Build Linux Kernel (2/5)

- ▶ `$ sudo apt-get install fakeroot build-essential kernel-package libncurses5 libncurses5-dev`
- ▶ `$ cd /usr/src`
- ▶ `$ sudo wget`
<https://cdn.kernel.org/pub/linux/kernel/v2.6/longterm/v2.6.32/linux-2.6.32.68.tar.xz>
- ▶ `$ sudo tar xvf linux-2.6.32.68.tar.xz`
- ▶ `$ cd linux-2.6.32.68`
- ▶ `$ sudo make mrproper`

Build Linux Kernel (3/5)

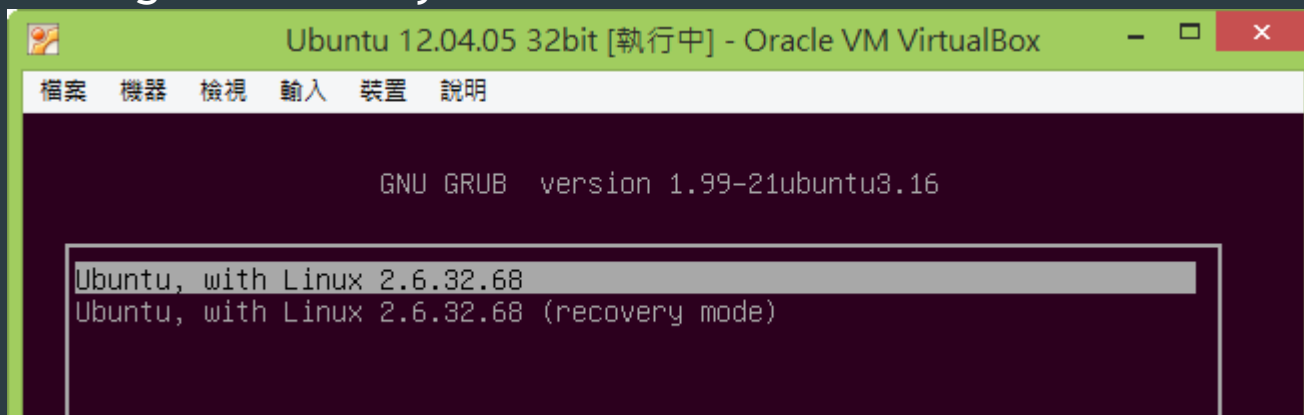
- ▶ `$ sudo make menuconfig`
- ▶ `$ sudo make bzImage`
 - ▶ You can use `make -j#` (# is the number of your physical cores) to create multiple threads to speed up the kernel building
- ▶ `$ sudo make modules`
- ▶ `$ sudo make modules_install`

Build Linux Kernel (4/5)

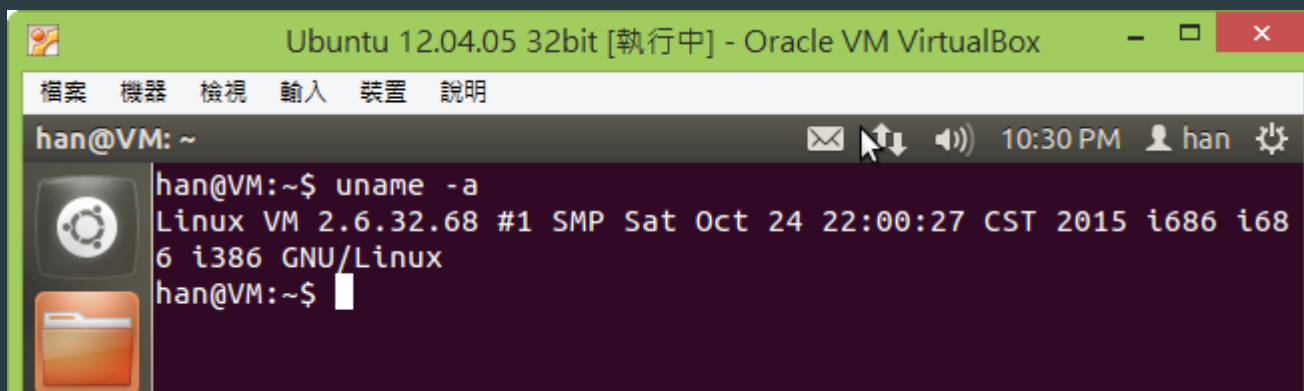
- ▶ `$ sudo make install`
- ▶ `$ sudo vim /etc/default/grub`
 - ▶ Add “#” to comment the following 2 lines
 - ▶ `#GRUB_HIDDEN_TIMEOUT=10`
 - ▶ `#GRUB_HIDDEN_TIMEOUT_QUIET=true`
- ▶ `$ sudo update-grub2`
- ▶ `$ sudo shutdown -r now`

Build Linux Kernel (5/5)

- Now, you can select the version 2.6.32.68 kernel in the GNU grub to boot your Ubuntu.



- Then, you can use terminal and type “uname -a” to check the kernel version.



References

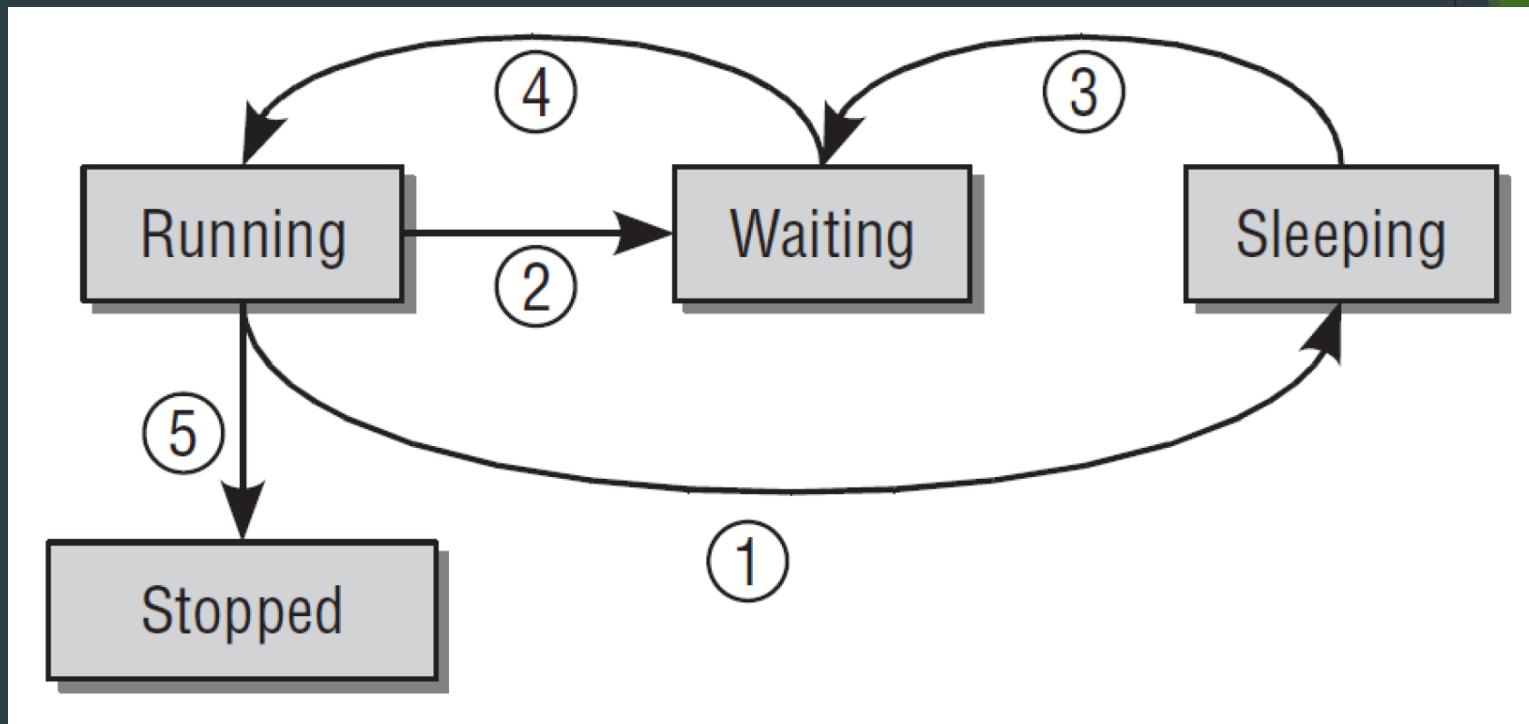
- ▶ [1] Wikipedia
[http://en.wikipedia.org/wiki/Kernel_\(computing\)](http://en.wikipedia.org/wiki/Kernel_(computing))
- ▶ [2] Oracle VM VirtualBox <https://www.virtualbox.org/>
- ▶ [3] Ubuntu <http://www.ubuntu.com/>
- ▶ [4] Linux Kernel in a Nutshell
<http://www.kroah.com/lkn/>

SCHEDULING IN LINUX (補充資料)

Process Life Cycle

- ▶ A process is not always ready to run.
- ▶ The scheduler must know the status of every process in the system when switching between tasks.
- ▶ A process may have one of the following **states**:
 - ▶ Running — The process is **executing at the moment**.
 - ▶ Waiting — The process is able to run but is not allowed to **because the CPU is allocated to another process**. The scheduler can select the process at the next task switch.
 - ▶ Sleeping — The process is sleeping and cannot run because it is **waiting for an external event**. The scheduler cannot select the process at the next task switch.
- ▶ The system saves all processes in a process table.

Transitions between Process States



The Need of the Scheduler

- ▶ A unique description of each process is held in memory and is linked with other processes by means of several structures.
- ▶ This is the situation facing the scheduler, whose task is *to share CPU time between the programs to create the illusion of concurrent execution*.
- ▶ This task is split into two different parts —
 - ▶ One relating to the **scheduling policy** and
 - ▶ The other to **context switching**

Scheduling in Linux (1/2)

- ▶ The schedule function is the starting point to an understanding of scheduling operations.
- ▶ It is defined in “`kernel/sched.c`” and is one of the most frequently invoked functions in the kernel code.
- ▶ Not only priority scheduling but also two other soft real-time policies required by the POSIX standard are implemented.
- ▶ E.g., completely fair scheduling, real-time scheduling and scheduling of the idle task, etc.

Scheduling in Linux (2/2)

- ▶ The scheduler uses a series of data structures to sort and manage the processes in the system.
- ▶ Scheduling can be activated in two ways:
 - ▶ **Main scheduler**: Either directly if a task goes to sleep or wants to yield the CPU for other reasons,
 - ▶ **Periodic scheduler**: Or by a periodic mechanism that is run with constant frequency to check from time to time if switching tasks is necessary
- ▶ Generic scheduler = Main + Periodic schedulers

Generic Scheduler

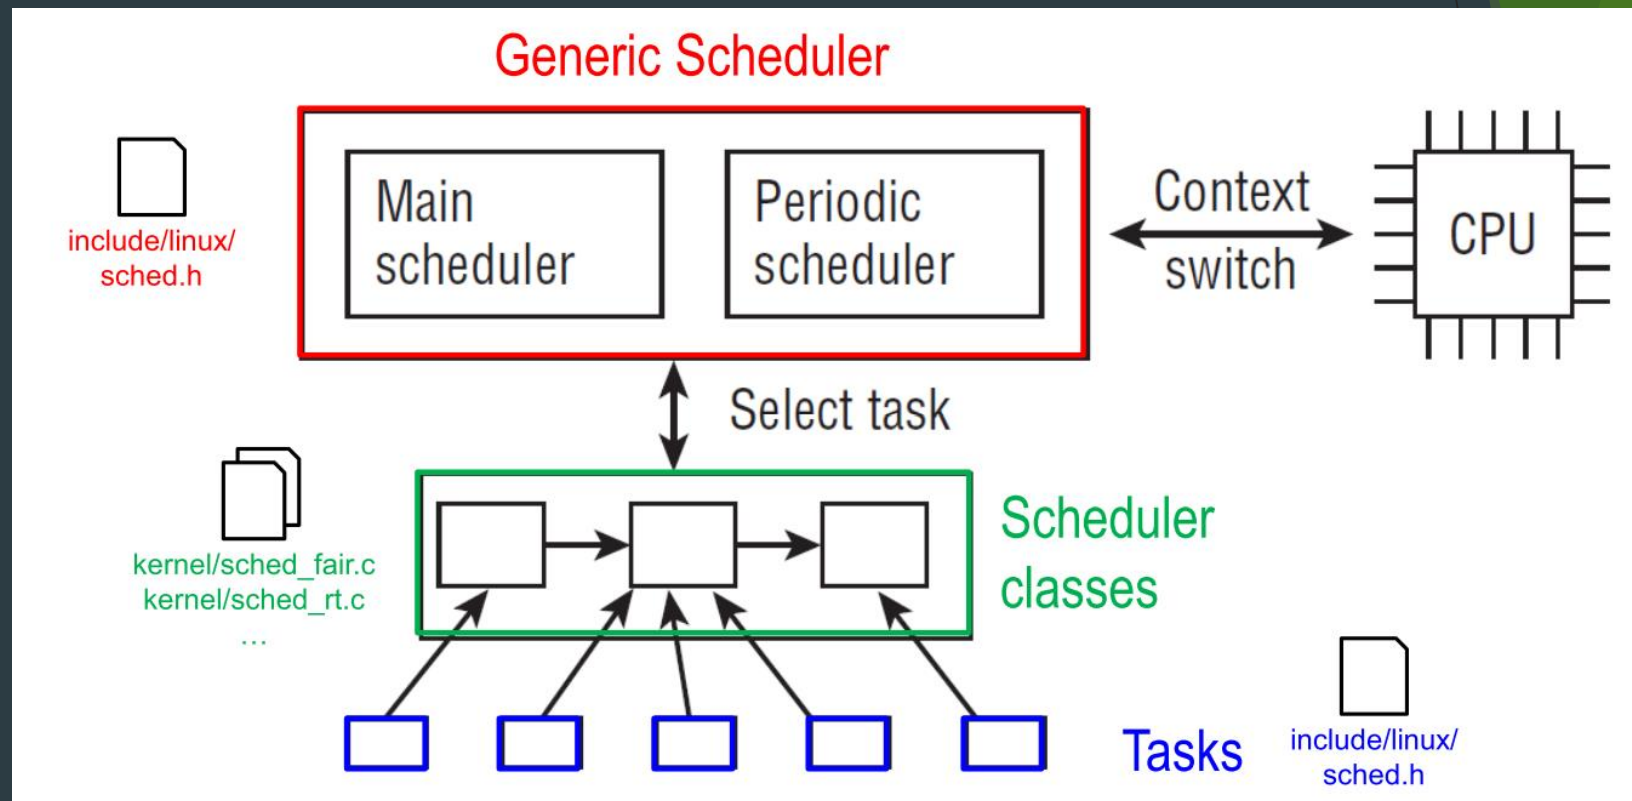
Scheduler Classes

Task

Task

Task

Linux Scheduling Subsystem



Linux Scheduling Policy Testing

Linux Scheduling Policy Testing

- ▶ Linux Scheduling Policy Classes
 - ▶ Normal Scheduling policies (Non-real-time)
 - ▶ SCHED_OTHER, SCHED_BATCH, SCHED_IDLE.
 - ▶ Real-Time policies
 - ▶ SCHED_FIFO, SCHED_RR.
- ▶ The default scheduling policy is **non-real-time**.
- ▶ In this part, using Linux **real-time scheduling policy (FIFO)** to schedule threads in a process.

Implement

- ▶ Write a C program (sched_test.c) to create two threads.
- ▶ Each thread will print who is running and busy for 1 second.
- ▶ Run the program by default time-sharing schedule policy and show the result.
Ex. `$./sched_test`
- ▶ Run the program by real-time scheduling policy (FIFO) and show the result.
Ex. `$./sched_test SCHED_FIFO`

```
1 int main(){
2
3     create thread 1
4     print "Thread 1 was created"
5
6     create thread 2
7     print "Thread 2 was created"
8
9 }
10
11 void thread_function(){
12
13     print "Thread # is running"
14     busy 1 second
15
16 }
```

Result

```
Terminal
root@han-X86: /home/han

root@han-X86:/home/han# ./sched_test
Thread 1 was created.
Thread 2 was created.
Thread 2 is running.
Thread 1 is running.
Thread 2 is running.
Thread 1 is running.
Thread 1 is running.
Thread 2 is running.
root@han-X86:/home/han#
root@han-X86:/home/han# ./sched_test SCHED_FIFO
Thread 1 was created.
Thread 2 was created.
Thread 1 is running.
Thread 1 is running.
Thread 1 is running.
Thread 2 is running.
Thread 2 is running.
Thread 2 is running.
root@han-X86:/home/han#
```

Hint

- ▶ `int sched_setscheduler(pid_t pid, int policy, const struct sched_param *param);`
- ▶ The policy corresponding value define in `/include/linux/sched.h`
- ▶ Set the priority of real-time process (`sched_param *param`)
- ▶ The permission to run real-time process
- ▶ The number of CPU cores on your virtual machine
 - ▶ CPU affinity

Project1 Requirements

- ▶ The project should
 - ▶ Contain **your report** (PDF format, within 2 pages) and **a C program**.
 - ▶ 報告內容可以是：編譯時遇到的問題、編譯步驟的解釋、關於 Linux Scheduling Policy Testing 與 Linux scheduling trace code 的心得、等等。
 - ▶ Please show your **name, student ID and E-mail** in your report
 - ▶ Be packed as one file named “**RTS_PJ1_Team##.zip**” or “**RTS_PJ1_StudentID.zip**” (If one member)
 - ▶ **RTS_PJ1_Team##.zip** or (**RTS_PJ1_StudentID.zip**)
 - **RTS_PJ1_report.pdf**
 - **sched_test.c**

Submission Rules

- ▶ Project deadline: 2015/11/09 (Monday) 23:59
 - ▶ Delayed submissions yield severe point deduction
- ▶ Send your project to TA
 - ▶ TA's E-mal: d03922006@csie.ntu.edu.tw
 - ▶ Mail title: RTS_PJ1_Team##
or “RTS_PJ1_StudentID” (If one member)
 - ▶ If you have new version: RTS_PJ1_Team##v#
or “RTS_PJ1_StudentIDv#” (If one member)
- ▶ DO NOT COPY THE HOMEWORK

References

- ▶ Reference Book

- ▶ Professional Linux® Kernel Architecture, Wolfgang Mauerer, Wiley Publishing, Inc.